

Programming Project #1: Hybrid Images

CS445: Computational Photography - Spring 2020

Part I: Hybrid Images

```
In [1]: import cv2

import numpy as np
from matplotlib.colors import LogNorm
from scipy import signal

import utils
```

```
In [2]: %matplotlib notebook
import matplotlib.pyplot as plt
```

```
In [3]: im1_file = './happy.jpg'
im2_file = './sad.jpg'

im1 = cv2.imread(im1_file, cv2.IMREAD_GRAYSCALE)
im2 = cv2.imread(im2_file, cv2.IMREAD_GRAYSCALE)
```

```
In [4]: pts_im1 = utils.prompt_eye_selection(im1)
```



```
In [5]: pts_im2 = utils.prompt_eye_selection(im2)
```



```
In [6]: im1, im2 = utils.align_images(im1_file, im2_file, pts_im1, pts_im2, save_in
```

```
In [7]: # convert to grayscale  
im1 = cv2.cvtColor(im1, cv2.COLOR_BGR2GRAY) / 255.0  
im2 = cv2.cvtColor(im2, cv2.COLOR_BGR2GRAY) / 255.0
```

```
In [8]: #Images sanity check
fig, axes = plt.subplots(1, 2)
axes[0].imshow(im1,cmap='gray')
axes[0].set_title('Image 1'), axes[0].set_xticks([]), axes[0].set_yticks([])
axes[1].imshow(im2,cmap='gray')
axes[1].set_title('Image 2'), axes[1].set_xticks([]), axes[1].set_yticks([])
```

Image 1



Image 2



```
In [9]: def hybridImage(im1, im2, cutoff_low, cutoff_high):
        ...
        Inputs:
            im1:    RGB (height x width x 3) or a grayscale (height x width)
                   as a numpy array.
            im2:    RGB (height x width x 3) or a grayscale (height x width)
                   as a numpy array.
            cutoff_low: standard deviation for the low-pass filter
            cutoff_high: standard deviation for the high-pass filter

        Output:
            Return the combination of both images, one filtered with a low-pass
            and the other with a high-pass filter.
        ...
        #https://stackoverflow.com/questions/47369579/how-to-get-the-gaussian

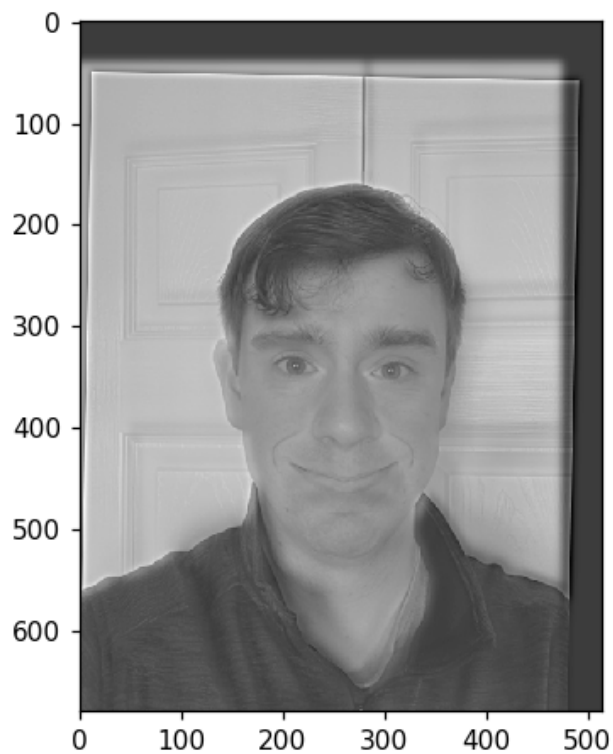
        sigma_low = 1./(2*np.pi*cutoff_low)
        sigma_high = 1./(2*np.pi*cutoff_high)

        fil_low = cv2.getGaussianKernel(25,sigma_low)
        fil_high = cv2.getGaussianKernel(25,sigma_high)
        fil_low = np.dot(fil_low,fil_low.T)
        fil_high = np.dot(fil_high,fil_high.T)
        im1_low = cv2.filter2D(im1, -1, fil_low)
        im2_low = cv2.filter2D(im2, -1, fil_high)
        return im2_low + im1 - im1_low, im1 - im1_low, im2_low
```

```
In [11]: #https://docs.opencv.org/2.4/modules/imgproc/doc/filtering.html#Mat%20ge
ksize_low = 20
ksize_high = 25
sigma_low = 0.3*((ksize_low-1)*0.5 - 1) + 0.8
sigma_high = 0.3*((ksize_high-1)*0.5 - 1) + 0.8
cutoff_low = 1./(2*np.pi*sigma_low)
cutoff_high = 1./(2*np.pi*sigma_high)

im_hybrid, im_high, im_low = hybridImage(im1, im2, cutoff_low, cutoff_hi

plt.close()
plt.imshow(im_hybrid, cmap='gray')
```



```
Out[11]: <matplotlib.image.AxesImage at 0x7ff9944832d0>
```


Below is the final cropped hybrid image. The foreground (up close) image is my smiling and the background (far away) image is me frowning. The foreground image contains only high frequency information and the background image contains only low frequencies. The hybrid image is created by simply adding these two images together.

```
In [12]: # Optional: Select top left corner and bottom right corner to crop image
# the function returns dictionary of
# {
#   'cropped_image': np.ndarray of shape H x W
#   'crop_bound': np.ndarray of shape 2x2
# }
cropped_object = utils.interactive_crop(im_hybrid)
```



This figure is an example of a missaligned mistake

```
In [13]: im_file = './mistake.png'  
im1 = cv2.imread(im_file, cv2.IMREAD_GRAYSCALE)  
plt.close()  
plt.imshow(im1, cmap='gray')
```




```
Out[13]: <matplotlib.image.AxesImage at 0x7ff994381b50>
```

These figures show the high frequency, low frequency, hybrid images and their cooresponding frequency plots.


```
In [14]: plt.close()
fig, axes = plt.subplots(2, 3)
axes[0,0].imshow(im_high,cmap='gray')
axes[0,0].set_title('High'), axes[0,0].set_xticks([]), axes[0,0].set_yticks([])
axes[0,1].imshow(im_low,cmap='gray')
axes[0,1].set_title('Low'), axes[0,1].set_xticks([]), axes[0,1].set_yticks([])
axes[0,2].imshow(im_hybrid,cmap='gray')
axes[0,2].set_title('Hybrid'), axes[0,2].set_xticks([]), axes[0,2].set_yticks([])

axes[1,0].imshow(np.log(np.abs(np.fft.fftshift(np.fft.fft2(im_high))))),cmap='gray')
axes[1,0].set_title('High FFT'), axes[1,0].set_xticks([]), axes[1,0].set_yticks([])
axes[1,1].imshow(np.log(np.abs(np.fft.fftshift(np.fft.fft2(im_low))))),cmap='gray')
axes[1,1].set_title('Low FFT'), axes[1,1].set_xticks([]), axes[1,1].set_yticks([])
axes[1,2].imshow(np.log(np.abs(np.fft.fftshift(np.fft.fft2(im_hybrid))))),cmap='gray')
axes[1,2].set_title('Hybrid FFT'), axes[1,2].set_xticks([]), axes[1,2].set_yticks([])
```



Color Hybrid Image (Bells and Whistles)

Changing emotion hybrid image example with color (bells and wistles). It seems best to use the color on the low frequency image. Using the color from the high frequency image seems to overwhelm the low frequencies.

```

In [15]: im1_file = './happy.jpg'
im2_file = './sad.jpg'

im1, im2 = utils.align_images(im1_file, im2_file, pts_im1, pts_im2, save_in

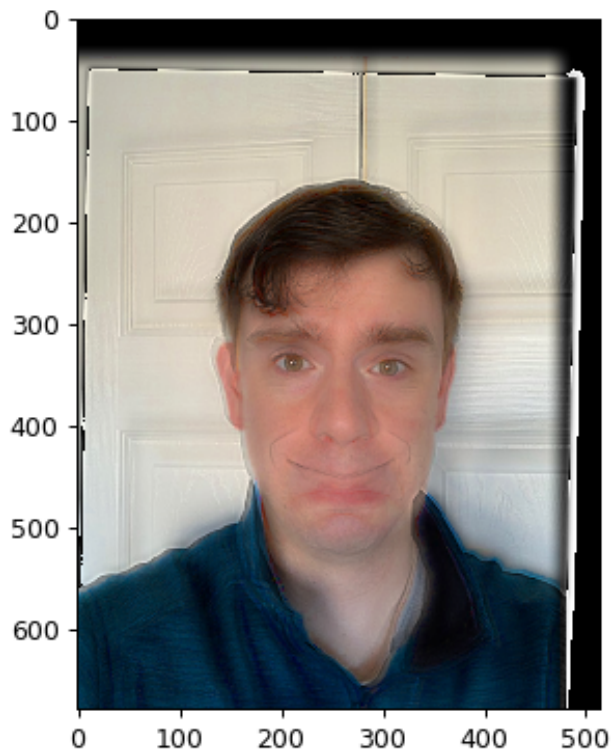
hsv1 = cv2.cvtColor(im1, cv2.COLOR_BGR2HSV)
hsv2 = cv2.cvtColor(im2, cv2.COLOR_BGR2HSV)
[h1, s1, v1] = [hsv1[:, :, 0], hsv1[:, :, 1], hsv1[:, :, 2]]
[h2, s2, v2] = [hsv2[:, :, 0], hsv2[:, :, 1], hsv2[:, :, 2]]

ksize_low = 15
ksize_high = 50
sigma_low = 0.3*((ksize_low-1)*0.5 - 1) + 0.8
sigma_high = 0.3*((ksize_high-1)*0.5 - 1) + 0.8
cutoff_low = 1./(2*np.pi*sigma_low)
cutoff_high = 1./(2*np.pi*sigma_high)

v_hybrid, v_high, v_low = hybridImage(v1, v2, cutoff_low, cutoff_high)

hsv_hyb = hsv2
hsv_hyb[:, :, 2] = v_hybrid
rgb = cv2.cvtColor(hsv_hyb, cv2.COLOR_HSV2RGB)
plt.close()
plt.imshow(rgb)

```



```

Out[15]: <matplotlib.image.AxesImage at 0x7ff9900b2350>

```

Another example of hybrid images fading a Tesla into a Porsche. Here the same

alignment tool was used to align the cars except that the center of the wheels were used for alignment rather than the center of the eyes.

In [16]:


```
im1_file = './tesla.jpg'  
im2_file = './porsche.jpg'  
  
im1 = cv2.imread(im1_file, cv2.IMREAD_GRAYSCALE)  
im2 = cv2.imread(im2_file, cv2.IMREAD_GRAYSCALE)
```

In [17]:

```
pts_im1 = utils.prompt_eye_selection(im1)
```



```
In [18]: pts_im2 = utils.prompt_eye_selection(im2)
```




Type *Markdown* and LaTeX: α^2

```
In [20]: im1, im2 = utils.align_images(im1_file, im2_file, pts_im1, pts_im2, save_in
im1 = cv2.cvtColor(im1, cv2.COLOR_BGR2GRAY) / 255.0
im2 = cv2.cvtColor(im2, cv2.COLOR_BGR2GRAY) / 255.0

#https://docs.opencv.org/2.4/modules/imgproc/doc/filtering.html#Mat%20ge
ksize_low = 75
ksize_high = 75
sigma_low = 0.3*((ksize_low-1)*0.5 - 1) + 0.8
sigma_high = 0.3*((ksize_high-1)*0.5 - 1) + 0.8
cutoff_low = 1./(2*np.pi*sigma_low)
cutoff_high = 1./(2*np.pi*sigma_high)
im_hybrid, im_high, im_low = hybridImage(im1, im2, cutoff_low, cutoff_hi

plt.close()
cropped_object = utils.interactive_crop(im_hybrid)
```



Type *Markdown* and LaTeX: α^2

Part II: Image Enhancement

Two out of three types of image enhancement are required. Choose a good image to showcase each type and implement a method. This code doesn't rely on the hybrid image part.

Contrast enhancement


```
In [21]: def hist_eq(im,alpha):  
    returnim = im  
    valuehist = np.histogram(im, bins=np.arange(256))  
    N = im.size  
    c = np.cumsum(valuehist[0])  
    for i in range(255):  
        returnim = np.where(returnim==i, alpha*c[i]*255./N + (1.-alpha)*  
            #np.put(returnim,np.argwhere(returnim==i),c[i]*255./N)  
    return returnim
```

Original Image:

```
In [22]: bc = './bad_contrast.jpg'

img = cv2.imread(bc)
hsv = cv2.cvtColor(img,cv2.COLOR_BGR2HSV)
rgb = cv2.cvtColor(hsv,cv2.COLOR_HSV2RGB)


plt.close()
plt.imshow(rgb)
```



```
Out[22]: <matplotlib.image.AxesImage at 0x7ff99298f0d0>
```

Contrast is corrected by performing histogram equalization on the value channel. Histogram equalization seemd to be the best approach to utilize the full range of intensities in the image.


```
In [23]: [h,s,v] = [hsv[:, :, 0], hsv[:, :, 1], hsv[:, :, 2]]  
  
veq = hist_eq(v, 0.8)  
hsv_eq = hsv  
hsv_eq[:, :, 2] = veq  
  
rgb = cv2.cvtColor(hsv_eq, cv2.COLOR_HSV2RGB)  
plt.close()  
plt.imshow(rgb)
```



```
Out[23]: <matplotlib.image.AxesImage at 0x7ff978fd8290>
```


Color enhancement

Original Image:

```
In [24]: bc = './bad_contrast.jpg'

img = cv2.imread(bc)
hsv = cv2.cvtColor(img,cv2.COLOR_BGR2HSV)
rgb = cv2.cvtColor(hsv,cv2.COLOR_HSV2RGB)


plt.close()
plt.imshow(rgb)
```



```
Out[24]: <matplotlib.image.AxesImage at 0x7ff978ff0510>
```

The colors are made more vivid by performing a histogram equalization on the saturation channel of the image.

```
In [25]: [h,s,v] = [hsv[:, :, 0], hsv[:, :, 1], hsv[:, :, 2]]  
  
sev = hist_eq(s, 0.5)  
hsv_eq = hsv  
hsv_eq[:, :, 1] = sev  
rgb = cv2.cvtColor(hsv_eq, cv2.COLOR_HSV2RGB)  
plt.close()  
plt.imshow(rgb)
```




```
Out[25]: <matplotlib.image.AxesImage at 0x7ff978f5a910>
```

Color shift (Bells and Whistles)

```
In [26]: def gamma(im,gamma):  
         returnim = im/255.  
         returnim = np.power(returnim,gamma)  
         return returnim*255
```

Original Image:

```
In [27]: bc = './bad_contrast.jpg'  
img = cv2.imread(bc)  
lab = cv2.cvtColor(img,cv2.COLOR_BGR2LAB)  
rgb = cv2.cvtColor(lab,cv2.COLOR_LAB2RGB)  
  
plt.close()  
plt.imshow(rgb)
```




```
Out[27]: <matplotlib.image.AxesImage at 0x7ff978ebdd50>
```

Make the image more red by applying a gamma adjustment on the a channel in the LAB colorspace.

```
In [28]: [l,a,b] = [lab[:, :, 0], lab[:, :, 1], lab[:, :, 2]]
          ag = gamma(a, .7)
          lab_shift = lab
          lab_shift[:, :, 1] = ag
          rgb = cv2.cvtColor(lab_shift, cv2.COLOR_LAB2RGB)

          plt.close()
          plt.imshow(rgb)
```




```
Out[28]: <matplotlib.image.AxesImage at 0x7ff978eb4390>
```

In this image the color is made less yellow by applying a negative gamma adjustment on the b channel in the LAB colorspace.


```
In [29]: lab = cv2.cvtColor(img,cv2.COLOR_BGR2LAB)
[l,a,b] = [lab[:,:,:0],lab[:,:,:1],lab[:,:,:2]]
bg = gamma(b, -0.5)
lab_shift = lab
lab_shift[:,:,:2] = bg
rgb = cv2.cvtColor(lab_shift,cv2.COLOR_LAB2RGB)

plt.close()
plt.imshow(rgb)
```

Figure 1



Reset origina

Out[29]: <matplotlib.image.AxesImage at 0x7ff978e18d90>

I belive I should receive all 100 points for the regular assignment and 15 points of bells and whistles. I completed all three image enhancement tasks (10 points) and the hybrid image in color (5 points).

